

ARPON KAPURIA

✉ arpkapuria@gmail.com | 🌐 Personal Website | 📞 +8801797288296 | 🇬🇧 Bangladesh

EDUCATION

National Institute of Technology Tiruchirappalli

Bachelor of Technology (Computer Science and Engineering)

Tiruchirappalli, India

November 2020 – June 2024

Kushtia Government College, Kushtia

Kushtia, Bangladesh

HSC (Science) | GPA: 5 / 5

2019

Kushtia Zilla College, Kushtia

Kushtia, Bangladesh

SSC (Science) | GPA: 5 / 5

2017

WORK EXPERIENCE

Advanced Machine Intelligence Research Lab

Dhaka, Bangladesh

Research Assistant | Advisor: Prof. M. Firoz Mridha

November 2025 – Present

- Engaged in research on Generative NLP and assisted in implementing deep learning models.

Indian Institute of Technology Bombay

Mumbai, India

Research Intern | MeDAL Lab | Advisor: Prof. Amit Sethi

May 2023 – July 2023

- Conducted an in-depth survey on self-supervised learning Frameworks and reproduced BYOL and simCLR in **PyTorch**. Experimented with loss functions and architectural variations to study their impact on learned representations.
- Implemented a DICOM anonymization pipeline using **Pydicom** to remove patient metadata from **5,000+ scans** and integrated CRAFT model to detect text regions in X-ray images for masking PHI.

National Institute of Technology Tiruchirappalli

Tiruchirappalli, India

Undergraduate R&D Assistant | Industrial Automation Lab

October 2022 – February 2023

Project: Assam Smart Home Automation

Advisors: Prof. M. Brindha, Prof. G. S. Ilango

- Developed a cross-platform application to automate energy operations and digital billing for a solar-powered microgrid in Assam, India, serving **20+ households**, as part of a govt. funded project.
- Integrated **Flutter** frontend with Django-based REST API and deployed an MQTT broker enabling real-time data exchange with **50+ IoT sensors** for energy consumption and control.

PROJECTS

Agentic RAG Pipeline with Hybrid Retrieval & Scalable Deployment

- Designed an end-to-end Agentic RAG system on with a **LangGraph** planner routing queries across hybrid retrieval (**Qdrant** vector search and **Neo4j** graph traversal), using HyDE, query rewriting, and a cross-encoder reranker to improve multi-hop reasoning beyond embedding-only pipelines.
- Deployed LLM inference via **vLLM** on **Ray Serve** with AWQ quantization, **continuous batching**, and **tensor parallelism**; served BGE-M3 embeddings with **fractional GPU allocation** and **torch.compile**, scaling to **128 concurrent sequences per replica** under autoscaled GPU worker groups.
- Provisioned multi-AZ AWS infrastructure with **Terraform** (**EKS**, **Aurora Serverless v2**, **ElastiCache HA**, **S3** with lifecycle tiering) and configured **Karpenter** with separate CPU and GPU provisioners – **30-second GPU scale-to-zero** and **spot consolidation** – reducing idle compute cost without compromising autoscaling responsiveness.
- Instrumented the full request path with **OpenTelemetry** traces, **Prometheus** metrics, and **Grafana** dashboards; evaluated retrieval and generation quality offline via **RAGAS** and an **LLM-as-a-Judge** pipeline against a curated golden dataset.

SKILLS

- Programming Languages:** Python, C, C++, SQL, Bash
- AI / ML:** PyTorch, Hugging Face Transformers, PEFT, TRL, Unsloth, LangChain, LangGraph, vLLM
- Backend & Data:** FastAPI, Ray (Data, Serve), PostgreSQL, Redis, Vector Databases
- Infrastructure & MLOps:** MLflow, DVC, Docker, Kubernetes, Terraform, AWS, Git, CI/CD
- Observability:** LangSmith, Prometheus, Grafana, OpenTelemetry

Last updated in June 2026